

Making Everything Easier!™

6th Edition

C++ FOR DUMMIES®

Learn to:

- Understand C++ 2009 programming from the ground up
- Write your first program in Chapter 1
- Master classes and inheritance
- Sail right through streaming I/O

Stephen R. Davis

Coauthor of C# 2008 For Dummies



Chapter 1

Writing Your First C++ Program

In This Chapter

- ▶ Finding out about C++
 - ▶ Installing Code::Blocks from the accompanying CD-ROM
 - ▶ Creating your first C++ program
 - ▶ Executing your program
-

Okay, so here we are: No one here but just you and me. Nothing left to do but get started. Might as well lay out a few fundamental concepts.

A computer is an amazingly fast but incredibly stupid machine. A computer can do anything you tell it (within reason), but it does *exactly* what it's told — nothing more and nothing less.

Perhaps unfortunately for us, computers don't understand any reasonable human language — they don't speak English either. Okay, I know what you're going to say: "I've seen computers that could understand English." What you really saw was a computer executing a *program* that could meaningfully understand English.

Computers understand a language variously known as *computer language* or *machine language*. It's possible but extremely difficult for humans to speak machine language. Therefore, computers and humans have agreed to sort of meet in the middle, using intermediate languages such as C++. Humans can speak C++ (sort of), and C++ can be converted into machine language for the computer to understand.

Grasping C++ Concepts

A C++ program is a text file containing a sequence of C++ commands put together according to the laws of C++ grammar. This text file is known as the *source file* (probably because it's the source of all frustration). A C++ source file normally carries the extension `.CPP` just as a Microsoft Word file ends in `.DOC` or an MS-DOS (remember that?) batch file ends in `.BAT`.

The point of programming in C++ is to write a sequence of commands that can be converted into a machine-language program that actually *does* what we want done. This is called *compiling* and is the job of the compiler. The machine code that you wrote must be combined with some setup and tear-down instructions and some standard library routines in a process known as *linking* or *building*. The resulting *machine-executable* files carry the extension .EXE in Windows.

That sounds easy enough — so what's the big deal? Keep going.

To write a program, you need two specialized computer programs. One (an editor) is what you use to write your code as you build your .CPP source file. The other (a compiler) converts your source file into a machine-executable .EXE file that carries out your real-world commands (open spreadsheet, make rude noises, deflect incoming asteroids, whatever).

Nowadays, tool developers generally combine compiler and editor into a single package — a development *environment*. After you finish entering the commands that make up your program, you need only click a button to create the executable file.

Fortunately, there are public-domain C++ environments. We use one of them in this book — the Code::Blocks environment. This editor will work with a lot of different compilers, but a version of Code::Blocks combined with the GNU gcc compiler for 32-bit versions of Windows is included on the book's CD-ROM. You can download the most recent version of Code::Blocks from www.codeblocks.org or you can download a recent version that's been tested for compatibility with the programs in this book from the author's Web site at www.stephendavis.com.

You can download versions of the gcc compiler for the Mac or Linux from www.gnu.org.

Although Code::Blocks is public domain, you're encouraged to pay some small fee to support its further development. You don't *have* to pay to use Code::Blocks, but you can contribute to the cause if you like. See the Web site for details.

I have tested the programs in this book with Code::Blocks combined with gcc version 4.4; the programs should work with later versions as well. You can check out my Web site for a list of any problems that may arise with future versions of Code::Blocks, gcc, or Windows.



Code::Blocks is a full-fledged editor and development environment front end. Code::Blocks supports a multitude of different compilers including the gcc compiler included on the enclosed CD-ROM.



The Code::Blocks/gcc package generates Windows-compatible 32-bit programs, but it does not easily support creating programs that have the classic Windows look. I strongly recommend that you work through the examples in this book first to learn C++ *before* you tackle Windows development. C++ and Windows programming are two separate things and (for the sake of your sanity) should remain so in your mind.

Follow the steps in the next section to install Code::Blocks and build your first C++ program. This program's task is to convert a temperature value entered by the user from degrees Celsius to degrees Fahrenheit.

Installing Code::Blocks

The CD-ROM that accompanies this book includes the most recent version of the Code::Blocks environment at the time of this writing.

The Code::Blocks environment comes in an easy-to-install, compressed executable file. This executable file is contained in the `CodeBlocks` directory on the accompanying CD-ROM. Here's the rundown on installing the environment:

1. **Insert the CD into the CD-ROM drive.**
2. **When the CD interface appears with the License Agreement, click Accept.**
3. **Click the Installing Code::Blocks button on the left.**
4. **Click Open Directory.**
5. **Double-click the `codeblocks_setup.exe` file.**

You can also choose Start→Run and type `x:\codeblocks\setup` in the window that appears, where *x* is the letter designation for your CD-ROM drive (normally D).

6. **Depending on what version of Windows you're using, you may get the ubiquitous "An unidentified programs wants access to your computer" warning pop-up. If so, click Allow to get the installation ball rolling.**
7. **Click Next after closing all extraneous applications as you are warned in the Welcome dialog to the CodeBlocks Setup Wizard.**
8. **Read the End User Legal Agreement (commonly known as the EULA) and then click I Agree if you can live with its provisions.**

It's not like you have much choice — the package really won't install itself if you don't accept. Assuming you *do* click OK, Code::Blocks opens a window showing the installation options. The default options are fine.

9. Click the Next button.

The installation program asks where you want it to install Code::Blocks. This dialog box also shows you how much disk space the installation requires (and whether you have enough). The default is okay assuming you have enough disk space (if not, you'll have to delete one of your reruns of the Simpsons).

10. Click Install.

Code::Blocks commences to copying a whole passel of files to your hard disk. Code::Blocks then asks "Do you want to run Code::Blocks now?"

11. Click Yes to start Code::Blocks.

Code::Blocks now asks which compiler you intend to use. The default is GNU GCC Compiler, which is the proper selection.

12. Click OK to select the GNU GCC compiler and start Code::Blocks.

Code::Blocks now wants to know which file associations you want to establish. The default is to allow Code::Blocks to open .CPP files. This option is fine unless you have another C++ compiler that you would rather have as the default.

13. Select the Yes, Associate C++ Files with Code::Blocks option if you do not have another C++ compiler installed. Otherwise, select the No, Leave Everything as It Is option. Click OK.

You now need to make sure that the compiler options are set to enable all warnings and C++ 2009 features.

14. From within Code::Blocks, choose Settings⇨Compiler and Debugger. In the Compiler Flags tab, make sure that the Enable All Compiler Warnings is selected.**15. Select the Enable All Compiler Warnings option, as shown in Figure 1-1.**

Start Code::Blocks. From within Code::Blocks, choose Settings⇨Compiler and debugger. In the Compiler Flags tab, make sure that the Enable All Compiler Warnings is selected.

16. Set the options to ensure C++ 2009 compliance.

The 2009 extensions are still considered a bit experimental as of this writing, so you need to tell gcc to enable these features. Click the Other Options tab and add the two lines `-std=c++0x` and `-Wc++0x-compat` as shown in Figure 1-2.

17. Click OK.**18. Click Next in the Code::Blocks Setup window and then click Finish to complete the setup program.**

The setup program exits.

Figure 1-1:
Ensure that
the Enable
All Compiler
Warnings
is set.

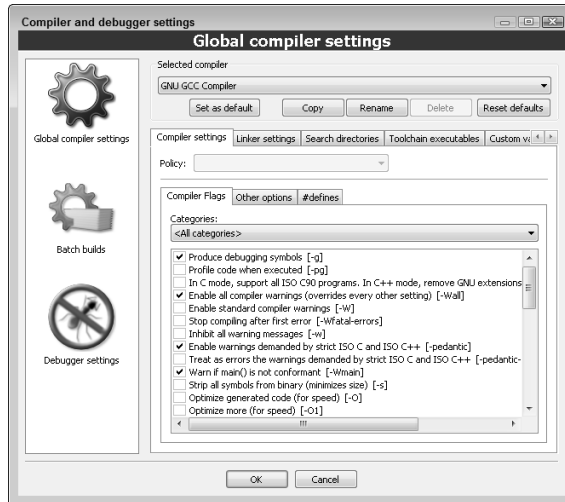
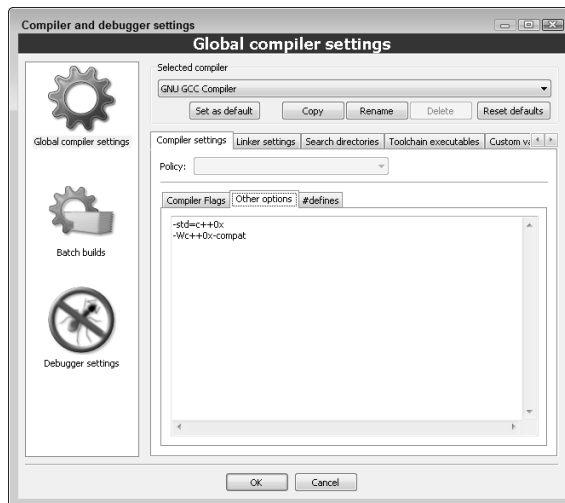


Figure 1-2:
Add these
lines to
enable the
C++ 2009
features.



Creating Your First C++ Program

In this section, you create your first C++ program. You enter the C++ code into a file called `CONVERT.CPP` and then convert the C++ code into an executable program.

Creating a project

The first step to creating a C++ program is to create what is known as a project. A *project* tells Code::Blocks the names of the .CPP source files to include and what type of program to create. Most of the programs in the book will consist of a single source file and will be command-line style:

1. Choose **Start**→**Programs**→**CodeBlocks**→**CodeBlocks to start up the CodeBlocks tool**.
2. From within Code::Blocks, choose **File**→**New**→**Project**.
3. Select the **Console Application** icon and then click **Go**.
4. Select **C++** as the language you want to use from the next window. Click **Next**.

Code::Blocks and gcc also support plain ol' C programs.

5. Select the **Console Application** icon and then click **Go**.
6. In the **Folder to Build Project In** field, navigate to the subdirectory where you want your program built.

I have divided the programs in this book by chapter, so I created the folder `C:\CPP_Programs\Chap01` using Windows Explorer and then selected it from the file menu.

7. In the **Project Title** field, type the name of the project, in this case **Conversion**.

The resulting screen is shown in Figure 1-3.

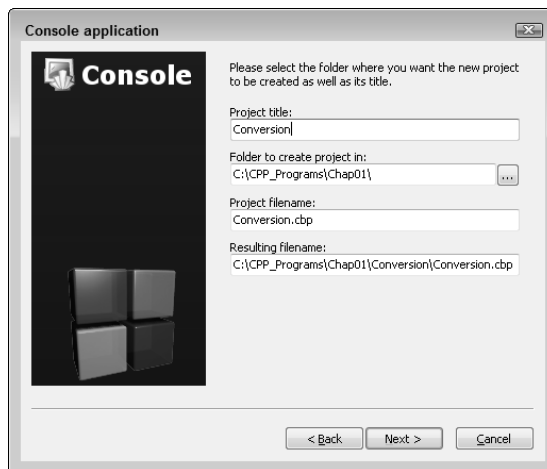


Figure 1-3:
I created
the project
Conversion
for the first
program.

8. Click Next.

The next window gives you the option of creating an application for testing or the final version. The default is fine.

9. Click Finish to create the Conversion project.

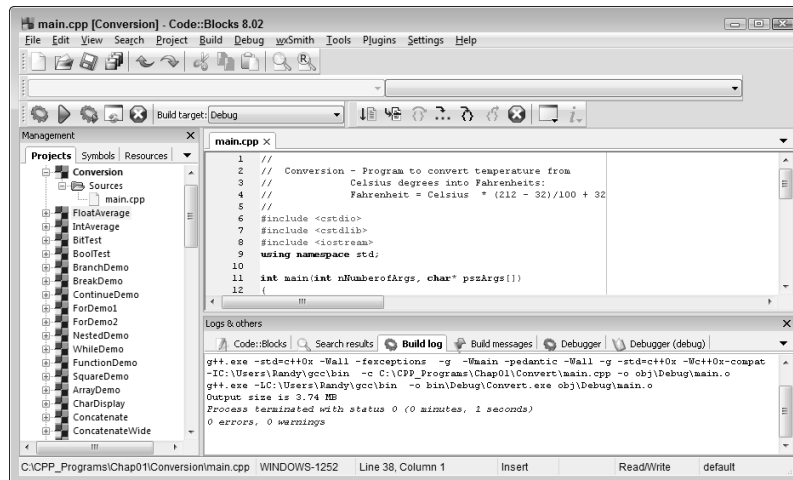
Entering the C++ code

The Conversion project that Code::Blocks creates consists of a single, default main.cpp file that does nothing. The next step is to enter our program:

1. In the Management window on the left, double-click main, which is under Sources, which is under Conversion.

Code::Blocks opens the empty main.cpp program that it created in the code editor, as shown in Figure 1-4. (Figure 1-4 shows all the projects from the entire book as well.)

Figure 1-4:
The Management window displays a directory structure for all available programs.

**2. Edit main.cpp with the following program exactly as written.**

Don't worry too much about indentation or spacing — it isn't critical whether a given line is indented two or three spaces, or whether there are one or two spaces between two words. C++ is case sensitive, however, so you need to make sure everything is lowercase.

You can cheat by using the files contained on the enclosed CD-ROM as described in the next section "Cheating."




```
//  
// Conversion - Program to convert temperature from  
//           Celsius degrees into Fahrenheit:  
//           Fahrenheit = Celsius * (212 - 32)/100 + 32  
//  
#include <stdio>  
#include <stdlib>  
#include <iostream>  
using namespace std;  
  
int main(int nNumberOfArgs, char* pszArgs[])  
{  
    // enter the temperature in Celsius  
    int celsius;  
    cout << "Enter the temperature in Celsius:";  
    cin >> celsius;  
  
    // calculate conversion factor for Celsius  
    // to Fahrenheit  
    int factor;  
    factor = 212 - 32;  
  
    // use conversion factor to convert Celsius  
    // into Fahrenheit values  
    int fahrenheit;  
    fahrenheit = factor * celsius/100 + 32;  
  
    // output the results (followed by a NewLine)  
    cout << "Fahrenheit value is:";  
    cout << fahrenheit << endl;  
  
    // wait until user is ready before terminating program  
    // to allow the user to see the program results  
    system("PAUSE");  
    return 0;  
}
```

3. Choose File→Save to save the source file.

I know that it may not seem all that exciting, but you've just created your first C++ program!

Cheating

All the programs in the book are included on the enclosed CD-ROM along with the project files to build them. You can use these files in two ways: one way is to go through all the steps to create the program by hand first but copy and paste from the sources on the CD-ROM into your program if you get into trouble (or your fingers start cramping). This is the preferred technique.

Alternatively you can use the following procedure to copy all the programs to your hard disk at once:

1. **Copy all the sources from the CD-ROM to the hard disk.**

This copies all the source code from the book along with the project files to build those programs.

2. **Double-click `AllPrograms.workspace` in `C:\CPP_Programs`.**

A workspace is a single file that references one or more projects. The `AllPrograms.workspace` file contains references to all the projects defined in the book.

3. **Right-click the `Conversion` project in the Management window on the left. Choose `Activate Project` from the context-sensitive menu that appears.**

Code::Blocks turns the `Conversion` label bold to verify that this is the program you are working with right now.

Building your program

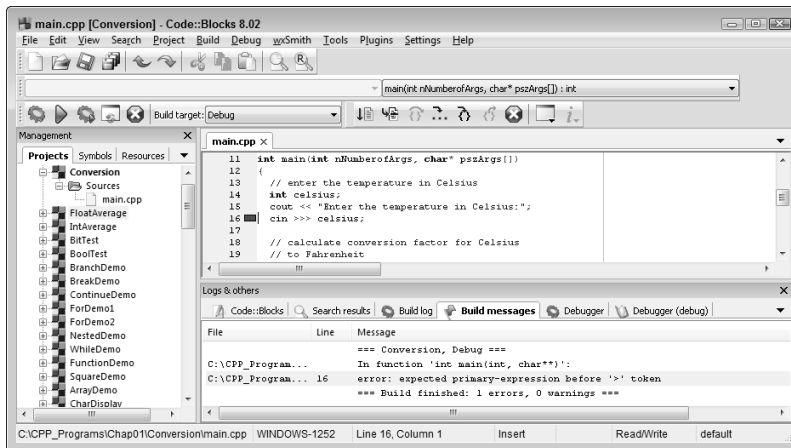
After you've saved your C++ source file to disk, it's time to generate the executable machine instructions.

To build your `Conversion` program, you choose `Build`⇨`Build` from the menu or press `Ctrl-F9`. Almost immediately, Code::Blocks takes off, compiling your program with gusto. If all goes well, the happy result of *0 errors, 0 warnings* appears in the lower-right window.

Code::Blocks generates a message if it finds any type of error in your C++ program — and coding errors are about as common as ice cubes in Alaska. You'll undoubtedly encounter numerous warnings and error messages, probably even when entering the simple `Conversion.cpp`. To demonstrate the error-reporting process, let's change Line 16 from `cin >> celsius;` to `cin >>> celsius;`.

This seems an innocent enough offense — forgivable to you and me perhaps, but not to C++. Choose `Build`⇨`Build` to start the compile and build process. Code::Blocks almost immediately places a red square next to the erroneous line as shown in Figure 1-5. The message in the Build Message tab is a rather cryptic error: `expected primary-expression before '>' token`. To get rid of the message, remove the extra `>` and recompile.

Figure 1-5:
Code::Blocks flags
the source
of errors
quickly.



You probably consider the error message generated by the example a little cryptic but give it time — you've been programming for only about 30 minutes now. Over time you'll come to understand the error messages generated by Code::Blocks and gcc much better.



Code::Blocks was able to point directly at the error this time but it isn't always that good. Sometimes it doesn't notice the error until the next line or the one after that, so if the line flagged with the error looks okay, start looking at its predecessor to see if the error is there.

Executing Your Program

It's now time to execute your new creation . . . that is, to run your program. You will run the `CONVERT.EXE` program file and give it input to see how well it works.

To execute the Conversion program, choose **Build** → **Build and Run** or press **F9**. This rebuilds the program if anything has changed and executes the program if the build is successful.

A window opens immediately, requesting a temperature in Celsius. Enter a known temperature, such as 100 degrees. After you press **Enter**, the program returns with the equivalent temperature of 212 degrees Fahrenheit as follows:

```
Enter the temperature in Celsius:100
Fahrenheit value is:212
Press any key to continue . . .
```

The message `Press any key to continue...` gives you the opportunity to read what you've entered before it goes away. Press Enter, and the window (along with its contents) disappears. Congratulations! You just entered, built, and executed your first C++ program.

Notice that Code::Blocks is not truly intended for developing Windows programs. In theory, you can write a Windows application by using Code::Blocks, but it isn't easy. (Building windowed applications is so much easier in Visual Studio.NET.)

Windows programs show the user a visually oriented output, all nicely arranged in onscreen windows. `Conversion.exe` is a 32-bit program that executes *under* Windows, but it's not a Windows program in the visual sense.

If you don't know what *32-bit program* means, don't worry about it. As I said, this book isn't about writing Windows programs. The C++ programs you write in this book have a *command line interface* executing within an MS-DOS box.

Budding Windows programmers shouldn't despair — you didn't waste your money. Learning C++ is a prerequisite to writing Windows programs. I think that they should be mastered separately: C++ first, Windows second.

Reviewing the Annotated Program

Entering data in someone else's program is about as exciting as watching someone else drive a car. You really need to get behind the wheel itself. Programs are a bit like cars as well. All cars are basically the same with small differences and additions — okay, French cars are a lot different than other cars, but the point is still valid. Cars follow the same basic pattern — steering wheel in front of you, seat below you, roof above you, and stuff like that.

Similarly, all C++ programs follow a common pattern. This pattern is already present in this very first program. We can review the `Conversion` program by looking for the elements that are common to all programs.

Examining the framework for all C++ programs

Every C++ program you write for this book uses the same basic framework, which looks a lot like this:

```
//  
//  Template - provides a template to be used as the  
//              starting point  
//  
// the following include files define the majority of  
// functions that any given program will need  
#include <cstdio>  
#include <cstdlib>  
#include <iostream>  
using namespace std;  
  
int main(int nNumberOfArgs, char* pszArgs[])  
{  
    // your C++ code starts here  
  
    // wait until user is ready before terminating program  
    // to allow the user to see the program results  
    system("PAUSE");  
    return 0;  
}
```

Without going into all the boring details, execution begins with the code contained in the open and closed braces immediately following the line beginning `main()`.

I've copied this code into a file called `Template.cpp` located in the `main CPP_Programs` folder on the enclosed CD-ROM.

Clarifying source code with comments

The first few lines in the `Conversion` program appear to be freeform text. Either this code was meant for human eyes or C++ is a lot smarter than I give it credit for. These first six lines are known as comments. *Comments* are the programmer's explanation of what she is doing or thinking when writing a particular code segment. The compiler ignores comments. Programmers (*good* programmers, anyway) don't.

A C++ comment begins with a double slash (`//`) and ends with a newline. You can put any character you want in a comment. A comment may be as long as you want, but it's customary to keep comment lines to no more than 80 characters across. Back in the old days — “old” is relative here — screens were limited to 80 characters in width. Some printers still default to 80 characters across when printing text. These days, keeping a single line to fewer than 80 characters is just a good practical idea (easier to read; less likely to cause eyestrain; the usual).

A newline was known as a *carriage return* back in the days of typewriters — when the act of entering characters into a machine was called *typing* and not *keyboarding*. A *newline* is the character that terminates a command line.



C++ allows a second form of comment in which everything appearing after a `/*` and before a `*/` is ignored; however, this form of comment isn't normally used in C++ anymore. (Later in this book, I describe the one case in which this type of comment is applied.)

It may seem odd to have a command in C++ (or any other programming language) that's specifically ignored by the computer. However, all computer languages have some version of the comment. It's critical that the programmer explain what was going through her mind when she wrote the code. A programmer's thoughts may not be obvious to the next colleague who tries to use or modify her program. In fact, the programmer herself may forget what her program meant if she looks at it months after writing the original code and has left no clue.

Basing programs on C++ statements

All C++ programs are based on what are known as C++ statements. This section reviews the statements that make up the program framework used by the Conversion program.

A statement is a single set of commands. Almost all C++ statements other than comments end in a semicolon. (You see one other exception in Chapter 10). Program execution begins with the first C++ statement after the open brace and continues through the listing, one statement at a time.

As you look through the program, you can see that spaces, tabs, and newlines appear throughout the program. In fact, I place a newline after every statement in this program. These characters are collectively known as *whitespace* because you can't see them on the monitor.



You may add whitespace anywhere you like in your program to enhance readability — except in the middle of a word:

```
See wha  
t I mean?
```

Although C++ may ignore whitespace, it doesn't ignore case. In fact, C++ is case sensitive to the point of obsession. The variable `fullspeed` and the variable `FullSpeed` have nothing to do with each other. The command `int` is completely understandable, but C++ has no idea what `INT` means. See what I mean about fast but stupid compilers?

Writing declarations

The line `int nCelsius;` is a declaration statement. A *declaration* is a statement that defines a variable. A *variable* is a “holding tank” for a value of some type. A variable contains a *value*, such as a number or a character.

The term *variable* stems from algebra formulas of the following type:

```
x = 10
y = 3 * x
```

In the second expression, `y` is set equal to 3 times `x`, but what is `x`? The variable `x` acts as a holding tank for a value. In this case, the value of `x` is 10, but we could have just as well set the value of `x` to 20 or 30 or `-1`. The second formula makes sense no matter what the value of `x` is.

In algebra, you’re allowed but not required to begin with a statement such as `x = 10`. In C++, the programmer must define the variable `x` before she can use it.

In C++, a variable has a type and a name. The variable defined on line 11 is called `celsius` and declared to hold an integer. (Why they couldn’t have just said *integer* instead of *int*, I’ll never know. It’s just one of those things you learn to live with.)

The name of a variable has no particular significance to C++. A variable must begin with the letters *A* through *Z*, the letters *a* through *z*, or an underscore (`_`). All subsequent characters must be a letter, a digit 0 through 9, or an underscore. Variable names can be as long as you want to make them.



It’s convention that variable names begin with a lowercase letter. Each new word *within* a variable begins with a capital letter, as in `myVariable`.



Try to make variable names short but descriptive. Avoid names such as `x` because `x` has no particular meaning. A variable name such as `lengthOfLineSegment` is much more descriptive.

Generating output

The lines beginning with `cout` and `cin` are known as input/output statements, often contracted to I/O statements. (Like all engineers, programmers love contractions and acronyms.)

The first I/O statement says “Output the phrase *Enter the temperature in Celsius to cout*” (pronounced “see-out”). `cout` is the name of the standard C++ output device. In this case, the standard C++ output device is your monitor.

The next line is exactly the opposite. It says, in effect, “Extract a value from the C++ input device and store it in the integer variable `celsius`.” The C++ input device is normally the keyboard. What we have here is the C++ analog to the algebra formula $x = 10$ just mentioned. For the remainder of the program, the value of `celsius` is whatever the user enters there.

Calculating Expressions

All but the most basic programs perform calculations of one type or another. In C++, an *expression* is a statement that performs a calculation. Said another way, an expression is a statement that *has a value*. An *operator* is a command that generates a value.

For example, in the Conversion example program — specifically in the two lines marked as a calculation expression — the program declares a variable *factor* and then assigns it the value resulting from a calculation. This particular command calculates the difference of 212 and 32; the operator is the minus sign (`-`), and the expression is `212-32`.

Storing the results of an expression

The spoken language can be very ambiguous. The term *equals* is one of those ambiguities. The word *equals* can mean that two things have the same value as in “a dollar equals one hundred cents.” Equals can also imply assignment, as in math when you say that “y equals 3 times x.”

To avoid ambiguity, C++ programmers call `=` the *assignment operator*, which says (in effect), “Store the results of the expression to the right of the `equal` sign in the variable to the left.” Programmers say that “`factor` is assigned the value 212 minus 32.” For short, you can say “`factor` gets 212 minus 32.”



Never say “`factor` is *equal* to 212 minus 32.” You’ll hear this from some lazy types, but you and I know better.

Examining the remainder of Conversion

The second expression in the Conversion program presents a slightly more complicated expression than the first. This expression uses the same mathematical symbols: `*` for multiplication, `/` for division, and `+` for addition. In this case, however, the calculation is performed on variables and not simply on constants.

The value contained in the variable called `factor` (which was calculated as the results of `212 - 32`, by the way) is multiplied by the value contained in `celsius` (which was input from the keyboard). The result is divided by 100 and summed with 32. The result of the total expression is assigned to the integer variable `fahrenheit`.

The final two commands output the string `Fahrenheit value is:` to the display, followed by the value of `fahrenheit` — and all so fast that the user scarcely knows it's going on.

Chapter 2

Declaring Variables Constantly

In This Chapter

- ▶ Declaring variables
- ▶ Declaring different types of variables
- ▶ Using floating-point variables
- ▶ Declaring and using other variable types

The most fundamental of all concepts in C++ is the *variable* — a variable is like a small box. You can store things in the box for later use, particularly numbers. The concept of a variable is borrowed from mathematics. A statement such as

```
x = 1
```

stores the value 1 in the variable `x`. From that point forward, the mathematician can use the variable `x` in place of the constant 1 — until he changes the value of `x` to something else.

Variables work the same way in C++. You can make the assignment

```
x = 1;
```

From that point forward in the execution of the program, until the value of `x` is changed, the *value of* `x` is 1. References to `x` are replaced by the value 1. In this chapter, you will find out how to declare and initialize variables in C++ programs. You will also see the different types of variables that C++ defines and when to use each.

Declaring Variables

A mathematician might write something like the following:

```
(x + 2) = y / 2  
x + 4 = y  
solve for x and y
```

Any reader who's had algebra realizes right off that the mathematician has introduced the variables x and y . But C++ isn't that smart. (Computers may be fast, but they're stupid.)

You have to announce each variable to C++ before you can use it. You have to say something soothing like this:

```
int x;  
x = 10;  
  
int y;  
y = 5;
```

These lines of code *declare* that a variable x exists, is of type `int`, and has the value 10; and that a variable y of type `int` also exists with the value 5. (The next section discusses variable types.) You can declare variables (almost) anywhere you want in your program — as long as you *declare the variable before you use it*.

Declaring Different Types of Variables

If you're on friendly terms with math (and who isn't?), you probably think of a variable in mathematics as an amorphous box capable of holding whatever you might choose to store in it. You might easily write something like the following:

```
x = 1  
x = 2.3  
x = "this is a sentence"
```

Alas, C++ is not that flexible. (On the other hand, C++ can do things that people can't do, such as add a billion numbers or so in a second, so let's not get too uppity.) To C++, there are different types of variables just as there are different types of storage bins. Some storage bins are so small that they can handle only a single number. It takes a larger bin to handle a sentence.



Some computer languages try harder to accommodate the programmer by allowing her to place different types of data in the same variable. These languages are called *weakly typed* languages. C++ is a *strongly typed* language — it requires the programmer to specifically declare each variable along with its exact type.

The variable type `int` is the C++ equivalent of an *integer* — a number that has no fractional part. (Integers are also known as *counting numbers* or *whole numbers*.)

Integers are great for most calculations. I made it through most of elementary school with integers. It isn't until I turned 11 or so that my teachers started mucking up the waters with fractions. The same is true in C++: More than 90 percent of all variables in C++ are declared to be of type `int`.

Unfortunately, `int` variables aren't adapted to every problem. For example, if you worked through the temperature-conversion program in Chapter 1, you might have noticed that the program has a potential problem — it can calculate temperatures to the nearest degree. No fractions of a degree are allowed. This integer limitation wouldn't affect daily use because it isn't likely that someone (other than a meteorologist) would get all excited about being off a fraction of a degree. There are plenty of cases, however, where this isn't the case — for example, you wouldn't want to come up a half mile short of the runway on your next airplane trip due to a navigational round-off.

Reviewing the limitations of integers in C++

The `int` variable type is the C++ version of an integer. `int` variables suffer the same limitations as their counting-number integer equivalents in math do.

Integer round-off

Lopping off the fractional part of a number is called *truncation*. Consider the problem of calculating the average of three numbers. Given three `int` variables — `nValue1`, `nValue2`, and `nValue3` — an equation for calculating the average is

```
int nAverage; int nValue1; int nValue2; int nValue3;  
nAverage = (nValue1 + nValue2 + nValue3) / 3;
```

Because all three values are integers, the sum is assumed to be an integer. Given the values 1, 2, and 2, the sum is 5. Divide that by 3, and you get $1\frac{2}{3}$, or 1.666. C++ uses slightly different rules: Given that all three variables `nValue1`, `nValue2`, and `nValue3` are integers, the sum is also assumed to be an integer. The result of the division of one integer by another integer is also an integer. Thus, the resulting value of `nAverage` is the unreasonable but logical value of 1.

The problem is much worse in the following mathematically equivalent formulation:

```
int nAverage; int nValue1; int nValue2; int nValue3;  
nAverage = nValue1/3 + nValue2/3 + nValue3/3;
```

Plugging in the same 1, 2, and 2 values, the resulting value of `nAverage` is 0 (talk about logical but unreasonable). To see how this can occur, consider that $\frac{1}{3}$ truncates to 0, $\frac{2}{3}$ truncates to 0, and $\frac{2}{3}$ truncates to 0. The sum of 0, 0, and 0 is 0. You can see that integer truncation can be completely unacceptable.

Limited range

A second problem with the `int` variable type is its limited range. A normal `int` variable can store a maximum value of 2,147,483,647 and a minimum value of -2,147,483,648 — roughly from positive 2 billion to negative 2 billion, for a total range of about 4 billion.



Two billion is a very large number: plenty big enough for most uses. But it's not large enough for some applications, including computer technology. In fact, your computer probably executes faster than 2 gigahertz, depending on how old your computer is. (*Giga* is the prefix meaning billion.) A single strand of communications fiber — the kind that's been strung back and forth from one end of the country to the other — can handle way more than 2 billion bits per second.

Solving the truncation problem

The limitations of `int` variables can be unacceptable in some applications. Fortunately, C++ understands decimal numbers that have a fractional part. (Mathematicians also call those *real numbers*.) Decimal numbers avoid many of the limitations of `int` type integers. To C++ all decimal numbers have a fractional part even if that fractional part is 0. In C++, the number 1.0 is just as much a decimal number as 1.5. The equivalent integer is written simply as 1. Decimal numbers can also be negative, such as -2.3.

When you declare variables in C++ that are decimal numbers, you identify them as floating-point or simply `float` variables. The term *floating-point* means the decimal point is allowed to float back and forth, identifying as many decimal places as necessary to express the value. Floating-point variables are declared in the same way as `int` variables:

```
float fValue1;
```

Once declared, you cannot change the type of a variable. `fValue1` is now a `float` and will be a `float` for the remainder of the program. To see how floating-point numbers fix the truncation problem inherent with integers, convert all the `int` variables to `float`. Here's what you get:

```
float fValue;  
fValue = 1.0/3.0 + 2.0/3.0 + 2.0/3.0;
```

is equivalent to

```
fValue = 0.333... + 0.666... + 0.666...;
```

which results in the value

```
fValue = 1.666...;
```



I have written the value `1.6666 . . .` as if the number of trailing 6s goes on forever. This is not necessarily the case. A `float` variable has a limit to the number of digits of accuracy, but it's a lot more than I can keep track of.



A constant that has a decimal point is assumed to be a floating-point value. However, the default type for a floating-point constant is something known as a double precision, which in C++ is called simply `double`, as we'll see in the next section.

The programs `IntAverage` and `FloatAverage` are available on the enclosed CD in the `CPP_Programs\Chap02` directory to demonstrate the round-off error inherent in integer variables.

Looking at the limits of floating-point numbers

Although floating-point variables can solve many calculation problems such as truncation, they have some limitations themselves — the reverse of those associated with integer variables. Floating-point variables can't be used to count things, are more difficult for the computer to handle, and also suffer from round-off error (though not nearly to the same degree as `int` variables).

Counting

You cannot use floating-point variables in applications where counting is important. This includes C++ constructs that count. C++ can't verify which whole number value is meant by a given floating-point number.

For example, it's clear to you and me that 1.0 is 1 but not so clear to C++. What about 0.9 or 1.1? Should these also be considered as 1? C++ simply avoids the problem by insisting on using `int` values when counting is involved.

Calculation speed

Historically, a computer processor can process integer arithmetic quicker than it can floating-point arithmetic. Thus, while a processor can add 1 million integer numbers in a given amount of time, the same processor may be able to perform only 200,000 floating-point calculations during the same period.

Calculation speed is becoming less of a problem as microprocessors get faster. In addition, today's general-purpose microprocessors include special floating-point circuitry on board to increase the performance of these operations. However, arithmetic on integer values is just a heck of a lot easier and faster than performing the same operation on floating-point values.

Loss of accuracy

Floating-point `float` variables have a precision of about 6 digits, and an extra-economy size, double-strength version of `float` known as a `double` can handle about 13 significant digits. This can cause round-off problems as well.

Consider that $\frac{1}{3}$ is expressed as 0.333 . . . in a continuing sequence. The concept of an infinite series makes sense in math but not to a computer because it has a finite accuracy. The `FloatAverage` program outputs 1.66667 as the average 1, 2, and 2 — that's a lot better than the 0 output by the `IntAverage` version but not even close to an infinite sequence.

C++ can correct for round-off error in a lot of cases. For example, on output, C++ can sometimes determine that the user really meant 1 instead of 0.999999. In other cases, even C++ cannot correct for round-off error.

Not-so-limited range

Although the `double` data type has a range much larger than that of an integer, it's still limited. The maximum value for an `int` is a skosh more than 2 billion. The maximum value of a `double` variable is roughly 10 to the 38th power. That's 1 followed by 38 zeroes; it eats 2 billion for breakfast. (It's even more than the national debt, at least at the time of this writing.)



Remember, only the first 13 digits or so have any meaning; the remaining 25 digits are noise having succumbed to floating-point round-off error.

Declaring Variable Types

So far in this chapter, I have been trumpeting that variables must be declared and that they must be assigned a type. Fortunately (ta-dah!), C++ provides a number of variable types. See Table 2-1 for a list of variables, their advantages, and limitations.

Table 2-1 Common C++ Variable Types		
<i>Variable</i>	<i>Defining a Constant</i>	<i>What It Is</i>
<code>int</code>	<code>1</code>	A simple counting number, either positive or negative.
<code>short int</code>	<code>---</code>	A potentially smaller version of <code>int</code> . It uses less memory but has a smaller range.
<code>long int</code>	<code>10L</code>	A potentially larger version of <code>int</code> . There is no difference between <code>long</code> and <code>int</code> with <code>gcc</code> .
<code>long long int</code>	<code>10LL</code>	A potentially even larger version of <code>int</code> .
<code>float</code>	<code>1.0F</code>	A single precision real number. This smaller version takes less memory than a <code>double</code> but has less accuracy and a smaller range.
<code>double</code>	<code>1.0</code>	A standard floating-point variable.
<code>long double</code>	<code>---</code>	A potentially larger floating-point number. On the PC, <code>long double</code> is used for the native size of the 80x86 floating-point processor, which is 80 bits.
<code>char</code>	<code>'c'</code>	A single <code>char</code> variable stores a single alphabetic or digital character. Not suitable for arithmetic.
<code>wchar_t</code>	<code>L'c'</code>	A larger character capable of storing symbols with larger character sets like Chinese.
<code>char string</code>	<code>"this is a string"</code>	A string of characters forms a sentence or phrase.
<code>bool</code>	<code>true</code>	The only other value is <code>false</code> . No, I mean, it's <i>really</i> false. Logically false. Not <i>false</i> as in fake or ersatz or . . . never mind.



The `long long int` and `long double` were officially introduced with C++ '09.

The integer types come in both signed and unsigned versions. Signed is always the default (for everything except `char` and `wchar_t`). The unsigned version is created by adding the keyword `unsigned` in front of the type in the declaration. The unsigned constants include a *U* or *u* in their type designation. Thus, the following declares an unsigned `int` variable and assigns it the value 10:

```
unsigned int uVariable;
uVariable = 10U;
```

The following statement declares the two variables `lVariable1` and `lVariable2` as type `long int` and sets them equal to the value 1, while `dVariable` is a double set to the value 1.0. Notice in the declaration of `lVariable2` that the `int` is assumed and can be left off:

```
// declare two long int variables and set them to 1
long int lVariable1
long lVariable2;      // int is assumed
lVariable1 = lVariable2 = 1;
// declare a variable of type double and set it to 1.0
double dVariable; dVariable = 1.0;
```



You can declare a variable and initialize it in the same statement:

```
int nVariable = 1;  // declare a variable and
                   // initialize it to 1
```

A `char` variable can hold a single character; a character string (which isn't really a variable type but works like one for most purposes) holds a string of characters. Thus, `'C'` is a `char` that contains the character `C`, whereas `"C"` is a string with one character in it. A rough analogy is that a `'C'` corresponds to a nail in your hand, whereas `"C"` corresponds to a nail gun with one nail left in the magazine. (Chapter 9 describes strings in detail.)



If an application requires a string, you've gotta provide one, even if the string contains only a single character. Providing nothing but the character just won't do the job.

Types of constants

A *constant value* is an explicit number or character (such as 1, 0.5, or `'c'`) that doesn't change. As with variables, every constant has a type. In an expression such as `n = 1`; the constant value 1 is an `int`. To make 1 a long integer, write the statement as `n = 1L`;. The analogy is as follows: 1 represents

a pickup truck with one ball in it, whereas `1L` is a dump truck also with one ball. The number of balls is the same in both cases, but the capacity of one of the containers is much larger.

Following the `int` to `long` comparison, `1.0` represents the value 1 but in a floating-point container. Notice, however, that the default for floating-point constants is `double`. Thus, `1.0` is a `double` number and not a `float`.

The constant values `true` and `false` are of type `bool`. In keeping with C++'s attention to case, `true` is a constant but `TRUE` has no meaning.

A variable can be declared constant when it is created via the keyword `const`:

```
const double PI = 3.14159; // declare a constant variable
```

A `const` variable must be initialized with a value when it is declared, and its value cannot be changed by any future statement.



Variables declared `const` don't have to be named with all capitals, but by convention they often are. This is just a hint to the reader that this so-called variable is, in fact, not.

I admit that it may seem odd to declare a variable and then say that it can't change. Why bother? Largely because carefully named `const` variables can make a program a lot easier to understand. Consider the following two equivalent expressions:

```
double dC = 6.28318 * dR; // what does this mean?
double dCircumference = 2 * PI * dRadius; // this is a
// lot easier to understand
```

It should be a lot clearer to the reader of this code that the second expression is multiplying the radius of something by 2π to calculate the circumference.

Range of Numeric Types

It may seem odd but the C++ standard doesn't say exactly how big a number each of the data types can accommodate. The standard speaks only to the relative size of each data type. For example, it says that the maximum `long int` is at least as large as the maximum `int`.

The authors of C++ weren't trying to be mysterious. They merely wanted to allow the compiler to implement the absolute fastest code possible for the base machine. The standard was designed to work for all different types of processors running different operating systems.

However, it is useful to know the limits for your particular implementation. Table 2-2 shows the size of each number type on a Windows PC using the Code::Blocks/gcc compiler that comes on the enclosed CD-ROM.

Table 2-2 Range of Numeric Types in Code::Block/gcc			
<i>Variable</i>	<i>Size (bytes)</i>	<i>Accuracy</i>	<i>Range</i>
short	2	exact	–32768 to 32767
int	4	exact	–2,147,483,648 to 2,147,483,647
long	4	exact	–2,147,483,648 to 2,147,483,647
long long int	8	exact	–9,223,372,036,854,775,808 to 9,223,372,036,854,775,807
float	4	7 digits	$\pm 3.4028 \times 10^{\pm 38}$
double	8	16 digits	$\pm 1.7977 \times 10^{\pm 308}$
long double	12	19 digits	$\pm 1.1897 \times 10^{\pm 4932}$

Attempting to calculate a number that's beyond the range of its type is known as an *overflow*. The C++ standard generally leaves the results of an overflow as undefined. That's another way that the definers of C++ remained flexible.



On the PC, a floating-point overflow results in an exception, which if not handled will cause your program to crash. (I don't discuss exceptions until Chapter 24.) As bad as that sounds, an integer overflow is worse — C++ silently generates an incorrect value without complaint.

C++ collision with filenames

Windows uses the backslash character to separate folder names in the path to a file. (This is a remnant of MS-DOS that Windows has not been able to shake.) Thus, `Root\FolderA\File` represents `File` within `FolderA`, which is a subdirectory of `Root`.

Unfortunately, MS-DOS's use of the backslash conflicts with the use of the backslash to indicate an escape character in C++. The character `\\` is a backslash in C++. The MS-DOS path `Root\FolderA\File` is represented in C++ as the string `"Root\\FolderA\\File"`.

Special characters

You can store any printable character you want in a `char` or `string` variable. You can also store a set of nonprintable characters that are used as character constants. See Table 2-3 for a description of these important nonprintable characters.

Table 2-3 Special Characters	
Character Constant	What It Is
'\n'	newline
'\t'	tab
'\040'	The character whose value is 40 in octal (see Chapter 4 for a discussion of number systems)
'\x20'	The character whose value is 20 in hexadecimal (this is the same as '\040')
'\0'	null (i.e., the character whose value is 0)
'\\'	backslash

You have already seen the newline character at the end of strings. This character breaks a string and puts the parts on separate lines. A newline character may appear anywhere within a string. For example:

```
"This is line 1\nThis is line 2"
```

appears on the output as

```
This is line 1
This is line 2
```

Similarly, the `\t` tab character moves output to the next tab position. (This position can vary, depending on the type of computer you're using to run the program.)

The numerical forms allow you to specify any nonprinting character that you like, but results may vary. The character represented by `0xFB`, for example, depends on the font and the character set (and may not be a legal character at all).

Because the backslash character is used to signify special characters, a character pair for the backslash itself is required. The character pair `\\` represents the backslash.

Wide Loads on Char Highway

The standard `char` variable is a scant 1 byte wide and can handle only 255 different characters. This is plenty enough for European languages but not big enough to handle symbol-based languages such as kanji.

Several standards have arisen to extend the character set to handle the demands of these languages. UTF-8 uses a mixture of 8-, 16-, and 32-bit characters to implement almost every kanji or hieroglyph you can think of but still remain compatible with simple 8-bit ASCII. UTF-16 uses a mixture of 16- and 32-bit characters to achieve an expanded character set, and UTF-32 uses 32 bits for all characters.



UTF stands for Unicode Transformation Format, from which it gets the common nickname Unicode.



Table 2-4 describes the different character types supported by C++. At first, C++ tried to get by with a vaguely defined wide character type, `wchar_t`. This type was intended to be the wide character type native to the application program's environment. C++ '09 introduced specific types for UTF-16 and UTF-32.



UTF-16 is the standard encoding for Windows and .NET applications. The `wchar_t` type refers to UTF-16 in the Code::Blocks/gcc compiler included on the CD-ROM.

Table 2-4 The C++ Character Types

<i>Variable</i>	<i>Example</i>	<i>What It Is</i>
<code>char</code>	<code>'c'</code>	ASCII character
<code>wchar_t</code>	<code>L'c'</code>	Character in wide format
<code>char_16t</code>	<code>u'c'</code>	UTF-16 character
<code>char_32t</code>	<code>U'c'</code>	UTF-32 character

Any of the character types in Table 2-4 can be combined into strings as well:

```
wchar_t* wideString = L"this is a wide string";
```

(Ignore the asterisk for now. I have a lot to say about its meaning in Chapter 8.)

Are These Calculations Really Logical?

C++ provides a logical variable called `bool`. The type `bool` comes from *Boole*, the last name of the inventor of the logical calculus. A Boolean variable has two values: `true` and `false`.



There are actually calculations that result in the value `bool`. For example, "`x` is equal to `y`" is either `true` or `false`.

Mixed Mode Expressions

C++ allows you to mix variable types in a single expression. That is, you are allowed to add an integer with a double precision floating-point value. In the following expression, for example, `nValue1` is allowed to be an `int`:

```
// in the following expression the value of nValue1
// is converted into a double before performing the
// assignment
int nValue1 = 1;
nValue1 + 1.0;
```

An expression in which the two operands are not the same type is called a *mixed-mode expression*. Mixed-mode expressions generate a value whose type is equal to the more capable of the two operands. In this case, `nValue1` is converted to a `double` before the calculation proceeds. Similarly, an expression of one type may be assigned to a variable of a different type, as in the following statement:

```
// in the following assignment, the whole
// number part of fVariable is stored into nVariable
double dVariable = 1.0;
int nVariable;
nVariable = dVariable;
```



You can lose precision or range if the variable on the left side of the assignment is smaller. In the preceding example, C++ truncates the value of `dVariable` before storing it in `nVariable`.

Converting a larger value type into a smaller value type is called *demotion*, whereas converting values in the opposite direction is known as *promotion*. Programmers say that the value of `int` variable `nVariable1` is promoted to a `double` in expressions such as the following:

```
int nVariable1 = 1;
double dVariable = nVariable1;
```

Naming conventions

You may have noticed that the name of each of the variables that I create begins with a special character that seems to have nothing to do with the name. These special characters are not special to C++ at all; they are merely meant

to jog the reader's memory and indicate the type of the variable. A partial list of these special characters follows. Using this convention, I can immediately recognize `dVariable` as a variable of type `double`, for example.

<i>Character</i>	<i>Type</i>
<code>n</code>	<code>int</code>
<code>l</code>	<code>long</code>
<code>f</code>	<code>float</code>
<code>d</code>	<code>double</code>
<code>c</code>	<code>character</code>
<code>sz</code>	<code>string</code>

Religious wars worse than the "Great Palin Debate of 2008" have broken out over whether or not this naming convention clarifies C++ code. It helps me, so I stick with it. Try it

for awhile. If after a few months, you don't think it helps, feel free to change your naming convention.



Mixed-mode expressions are not a good idea. Avoid forcing C++ to do your conversions for you.

Automatic Declarations



This entire section works only for C++ '09.

If you are really lazy, you can let C++ determine the types of your variables for you. Consider the following declaration:

```
int nVar = 1;
```

You might ask, "Why can't C++ figure out the type of `nVar`?" The answer is, as of C++ '09, it will if you ask nicely, as follows:

```
auto var1 = 1;  
auto var2 = 2.0;
```

This says, “declare `var1` to be a variable of the same type as the constant value 1 (which happens to be an `int`) and declare `var2` to be the same type as 2.0 (which is a `double`).”



I consider the term `auto` to be a particularly unfortunate choice for this purpose because prior to C++ '09, the keyword `auto` had a completely different meaning. However, `auto` had fallen out of use for at least ten years, so the standards people figured that it would be safe to usurp the term. Just be aware that if you see the keyword `auto` in some old code, you will need to remove it.

You can also tell C++ that you want a variable to be declared to be of the same type as another variable, whatever that might be, using the keyword `decltype()`.

```
int var1;  
decltype(var1) var2; // declare var2 to be of the  
                     // same type as var1
```

C++ replaces the `decltype(var1)` with the type of `var1`, again an `int`.

Contents at a Glance



<i>Introduction</i>	<i>1</i>
<i>Part I: Introduction to C++ Programming</i>	<i>7</i>
Chapter 1: Writing Your First C++ Program.....	9
Chapter 2: Declaring Variables Constantly.....	25
Chapter 3: Performing Mathematical Operations.....	41
Chapter 4: Performing Logical Operations	49
Chapter 5: Controlling Program Flow.....	63
<i>Part II: Becoming a Functional C++ Programmer</i>	<i>79</i>
Chapter 6: Creating Functions.....	81
Chapter 7: Storing Sequences in Arrays	95
Chapter 8: Taking a First Look at C++ Pointers.....	111
Chapter 9: Taking a Second Look at C++ Pointers.....	127
Chapter 10: The C++ Preprocessor.....	145
<i>Part III: Introduction to Classes</i>	<i>159</i>
Chapter 11: Examining Object-Oriented Programming.....	161
Chapter 12: Adding Class to C++.....	165
Chapter 13: Point and Stare at Objects.....	181
Chapter 14: Protecting Members: Do Not Disturb.....	197
Chapter 15: Why Do You Build Me Up, Just to Tear Me Down Baby?.....	205
Chapter 16: Making Constructive Arguments	217
Chapter 17: Copying the Copy Copy Copy Constructor	235
Chapter 18: Static Members: Can Fabric Softener Help?	247
<i>Part IV: Inheritance</i>	<i>255</i>
Chapter 19: Inheriting a Class	257
Chapter 20: Examining Virtual Member Functions: Are They for Real?	265
Chapter 21: Factoring Classes	273

<i>Part V: Optional Features</i>	281
Chapter 22: A New Assignment Operator, Should You Decide to Accept It	283
Chapter 23: Using Stream I/O	293
Chapter 24: Handling Errors — Exceptions.....	315
Chapter 25: Inheriting Multiple Inheritance	325
Chapter 26: Tempting C++ Templates	337
Chapter 27: Standardizing on the Standard Template Library	345
<i>Part VI: The Part of Tens</i>	357
Chapter 28: Ten Ways to Avoid Adding Bugs to Your Program	359
Chapter 29: Ten Major Recent Additions to C++	367
<i>Appendix: About the CD</i>	377
<i>Index</i>	383

Table of Contents

Introduction 1

What's in This Book	1
What's on the CD	2
What Is C++?	2
Conventions Used in This Book.....	3
How This Book Is Organized	3
And There's More	4
Part I: Introduction to C++ Programming	4
Part II: Becoming a Functional C++ Programmer.....	4
Part III: Introduction to Classes	5
Part IV: Inheritance.....	5
Part V: Optional Features	5
Part VI: The Part of Tens	5
Icons Used in This Book	6
Where to Go from Here.....	6

Part I: Introduction to C++ Programming 7

Chapter 1: Writing Your First C++ Program9

Grasping C++ Concepts.....	9
Installing Code::Blocks.....	11
Creating Your First C++ Program.....	13
Creating a project	14
Entering the C++ code	15
Cheating	16
Building your program	17
Executing Your Program.....	18
Reviewing the Annotated Program	19
Examining the framework for all C++ programs	19
Clarifying source code with comments	20
Basing programs on C++ statements.....	21
Writing declarations	22
Generating output	22
Calculating Expressions.....	23
Storing the results of an expression.....	23
Examining the remainder of Conversion	24



Chapter 2: Declaring Variables Constantly	25
Declaring Variables	25
Declaring Different Types of Variables	26
Reviewing the limitations of integers in C++	27
Solving the truncation problem	28
Looking at the limits of floating-point numbers	29
Declaring Variable Types	31
Types of constants	32
Range of Numeric Types	33
Special characters	35
Wide Loads on Char Highway	36
Are These Calculations Really Logical?	37
Mixed Mode Expressions	37
Automatic Declarations	38
Chapter 3: Performing Mathematical Operations	41
Performing Simple Binary Arithmetic	41
Decomposing Expressions	43
Determining the Order of Operations	43
Performing Unary Operations	44
Using Assignment Operators	46
Chapter 4: Performing Logical Operations	49
Why Mess with Logical Operations?	49
Using the Simple Logical Operators	50
Storing logical values	51
Using logical int variables	53
Be careful performing logical operations on floating-point variables	53
Expressing Binary Numbers	55
The decimal number system	55
Other number systems	56
The binary number system	56
Performing Bitwise Logical Operations	58
The single-bit operators	58
Using the bitwise operators	60
A simple test	60
Do something logical with logical calculations	62
Chapter 5: Controlling Program Flow	63
Controlling Program Flow with the Branch Commands	63
Executing Loops in a Program	65
Looping while a condition is true	66
Using the autoincrement/autodecrement feature	67
Using the for loop	69
Avoiding the dreaded infinite loop	72
Applying special loop controls	73

Nesting Control Commands	75
Switching to a Different Subject?	77

Part 11: Becoming a Functional C++ Programmer..... 79

Chapter 6: Creating Functions81

Writing and Using a Function.....	81
Defining our first function.....	84
Defining the sumSequence() function.....	84
Calling the function sumSequence()	85
Divide and conquer	85
Understanding the Details of Functions	85
Understanding simple functions.....	86
Understanding functions with arguments	87
Overloading Function Names.....	90
Defining Function Prototypes	91
Variable Storage Types.....	93

Chapter 7: Storing Sequences in Arrays95

Arraying the Arguments for Arrays.....	95
Using an array	97
Initializing an array	100
Accessing too far into an array	101
Using arrays.....	101
Defining and using arrays of arrays.....	101
Using Arrays of Characters	102
Creating an array of characters	102
Creating a string of characters.....	103
Manipulating Strings with Character	105
Adding Some Library Functions	107
Making Room for Wide Strings	109

Chapter 8: Taking a First Look at C++ Pointers.....111

Variable Size.....	111
What's in an Address?.....	113
Address Operators	113
Using Pointer Variables	115
Using different types of pointers	116
Passing Pointers to Functions.....	117
Passing by value.....	118
Passing pointer values	118
Passing by reference	119
Constant const Irritation	120
Making Use of a Block of Memory Called the Heap.....	122
Limited scope	122
Examining the scope problem.....	123
Providing a solution using the heap.....	124

Chapter 9: Taking a Second Look at C++ Pointers	127
Defining Operations on Pointer Variables	127
Reexamining arrays in light of pointer variables	128
Applying operators to the address of an array	130
Expanding pointer operations to a string	131
Justifying pointer-based string manipulation	133
Applying operators to pointer types other than char	134
Contrasting a pointer with an array	134
When Is a Pointer Not?	136
Declaring and Using Arrays of Pointers	137
Utilizing arrays of character strings	138
Accessing the arguments to main()	140
Chapter 10: The C++ Preprocessor	145
What Is a Preprocessor?	145
Including Files	146
#Defining Things	149
Okay, how about not #defining things?	152
Enumerating other options	153
Including Things #if I Say So	154
Intrinsically Defined Objects	156
Typedef	158
 Part III: Introduction to Classes	 159
Chapter 11: Examining Object-Oriented Programming	161
Abstracting Microwave Ovens	161
Preparing functional nachos	162
Preparing object-oriented nachos	163
Classifying Microwave Ovens	163
Why Classify?	164
Chapter 12: Adding Class to C++	165
Introducing the Class	165
The Format of a Class	166
Accessing the Members of a Class	167
Activating Our Objects	167
Simulating real-world objects	168
Why bother with member functions?	168
Adding a Member Function	169
Calling a Member Function	170
Accessing other members from a member function	172
Scope Resolution (And I Don't Mean How Well Your Microscope Works)	173
Defining a Member Function in the Class	175

Keeping a Member Function After Class	176
Overloading Member Functions	178
Chapter 13: Point and Stare at Objects.	181
Declaring Arrays of Objects	181
Declaring Pointers to Objects	182
Dereferencing an object pointer	183
Pointing toward arrow pointers.....	184
Passing Objects to Functions.....	184
Calling a function with an object value	185
Calling a function with an object pointer	186
Calling a function by using the reference operator	188
Why Bother with Pointers or References?	189
Returning to the Heap	189
Allocating heaps of objects	190
Comparing Pointers to References.....	191
Linking Up with Linked Lists	191
Performing other operations on a linked list	192
Hooking up with a LinkedListData program.....	193
A Ray of Hope: A List of Containers Linked to the C++ Library	196
Chapter 14: Protecting Members: Do Not Disturb	197
Protecting Members.....	197
Why you need protected members	197
Discovering how protected members work	198
Making an Argument for Using Protected Members.....	200
Protecting the internal state of the class.....	200
Using a class with a limited interface.....	201
Giving Nonmember Functions Access to Protected Members	201
Chapter 15: Why Do You Build Me Up, Just to Tear Me Down Baby?	205
Creating Objects	205
Using Constructors.....	206
Constructing a single object.....	207
Constructing multiple objects.....	208
Constructing a duplex.....	209
Dissecting a Destructor	211
Why you need the destructor	211
Working with destructors.....	212
Chapter 16: Making Constructive Arguments	217
Outfitting Constructors with Arguments.....	217
Using a constructor	218
Placing Too Many Demands on the Carpenter:	
Overloading the Constructor	220
Defaulting Default Constructors	222

Constructing Class Members	224
Constructing a complex data member	224
Constructing a constant data member	228
Reconstructing the Order of Construction	229
Local objects construct in order	229
Static objects construct only once	230
All global objects construct before main()	231
Global objects construct in no particular order	231
Members construct in the order in which they are declared	233
Destructors destruct in the reverse order of the constructors	233
Constructors as a Form of Conversion	233
Chapter 17: Copying the Copy Copy Copy Constructor	235
Copying an Object	235
Why you need the copy constructor	236
Using the copy constructor	236
The Automatic Copy Constructor	238
Creating Shallow Copies versus Deep Copies	240
It's a Long Way to Temporaries	244
Avoiding temporaries, permanently	246
Chapter 18: Static Members: Can Fabric Softener Help?	247
Defining a Static Member	247
Why you need static members	247
Using static members	248
Referencing static data members	249
Uses for static data members	250
Declaring Static Member Functions	251
What Is This About Anyway?	253
Part IV: Inheritance	255
Chapter 19: Inheriting a Class	257
Do I Need My Inheritance?	258
How Does a Class Inherit?	259
Using a subclass	261
Constructing a subclass	262
Destructing a subclass	263
Having a HAS_A Relationship	263
Chapter 20: Examining Virtual Member Functions: Are They for Real?	265
Why You Need Polymorphism	268
How Polymorphism Works	269
When Is a Virtual Function Not?	270
Considering Virtual Considerations	271

Chapter 21: Factoring Classes	273
Factoring.....	273
Implementing Abstract Classes	277
Describing the abstract class concept.....	278
Making an honest class out of an abstract class	279
Passing abstract classes	280
 Part V: Optional Features	 281
 Chapter 22: A New Assignment Operator, Should You Decide to Accept It.	 283
Comparing Operators with Functions	283
Inserting a New Operator	284
Creating Shallow Copies Is a Deep Problem	285
Overloading the Assignment Operator	286
Overloading the Subscript Operator	291
 Chapter 23: Using Stream I/O	 293
How Stream I/O Works.....	293
Default stream objects	294
Stream Input/Output	295
Open modes.....	297
Hey, file, what state are you in?	298
Can you show me an example?	298
Other Methods of the Stream Classes	301
Reading and writing streams directly	303
Controlling format	305
What's up with endl?	307
Positioning the pointer within a file	307
Using the stringstream Subclasses	308
Manipulating Manipulators	311
 Chapter 24: Handling Errors — Exceptions.	 315
Justifying a New Error Mechanism?	317
Examining the Exception Mechanism	318
What Kinds of Things Can I Throw?	321
Just Passing Through.....	324
 Chapter 25: Inheriting Multiple Inheritance	 325
Describing the Multiple Inheritance Mechanism	325
Straightening Out Inheritance Ambiguities.....	327
Adding Virtual Inheritance	328
Constructing the Objects of Multiple Inheritance.....	335
Voicing a Contrary Opinion.....	335

Chapter 26: Tempting C++ Templates 337

Generalizing a Function into a Template.....	338
Template Classes.....	340
Tips for Using Templates.....	343

Chapter 27: Standardizing on the Standard Template Library 345

The string Container	346
Iterating through Lists	351
Making your way through a list	353
Operations on an entire list.....	354
Can you show me an example?	355

Part VI: The Part of Tens*..... 357*Chapter 28: Ten Ways to Avoid Adding Bugs to Your Program 359**

Enable All Warnings and Error Messages.....	359
Adopt a Clear and Consistent Coding Style	360
Limit the Visibility	360
Comment Your Code While You Write It.....	362
Single-Step Every Path at Least Once	362
Avoid Overloading Operators.....	363
Manage the Heap Systematically.....	363
Use Exceptions to Handle Errors.....	363
Declare Destructors Virtual	364
Avoid Multiple Inheritance.....	365

Chapter 29: Ten Major Recent Additions to C++. 367

Use Smart Pointers	367
Initialize Variables with a Variable-Length List	368
Initialize Data Members Inline	369
Instantiate an Extern Template.....	370
Implement Thread Local Storage	371
Use Rvalue References.....	371
Implement Concepts	372
Define Lambda Expressions	373
Define Variadic Templates	374
Use typeid().....	374

<i>Appendix: About the CD</i>	377
System Requirements	377
Using the CD	378
What You'll Find on the CD	378
CPP programs.....	379
Code::Blocks development environment.....	379
Troubleshooting	380
Customer Care	381
 <i>Index</i>	 383

